

# SGD Algorithm to predict movie ratings

There will be some functions that start with the word "grader" ex: grader\_matrix(), grader\_mean(), grader\_dim() etc, you should not change those function definition.

Every Grader function has to return True.

1. Download the data from [here](#)
2. The data will be of this format, each data point is represented as a triplet of user\_id, movie\_id and rating

user_id	movie_id	rating
77	236	3
471	208	5
641	401	4
31	298	4
58	504	5
235	727	5

## Task 1

### Predict the rating for a given (user\_id, movie\_id) pair


Predicted rating  $\hat{y}_{ij}$  for user  $i$ , movie  $j$  pair is calculated as  $\hat{y}_{ij} = \mu + b_i + c_j + u_i^T v_j$ , here we will be finding the best values of  $b_i$  and  $c_j$  using SGD algorithm with the optimization problem for  $N$  users and  $M$  movies is defined as

$$L = \min_{b, c, \{u_i\}_{i=1}^N, \{v_j\}_{j=1}^M} \alpha \left( \sum_j \sum_k v_{jk}^2 + \sum_i \sum_k u_{ik}^2 + \sum_i b_i^2 + \sum_j c_j^2 \right) + \sum_{i,j \in \mathcal{I}^{\text{train}}} (y_{ij} - \mu - b_i - c_j - u_i^T v_j)^2$$

- $\mu$  : scalar mean rating
- $b_i$  : scalar bias term for user  $i$
- $c_j$  : scalar bias term for movie  $j$
- $u_i$  : K-dimensional vector for user  $i$
- $v_j$  : K-dimensional vector for movie  $j$

- \*. We will be giving you some functions, please write code in that functions only.
- \*. After every function, we will be giving you expected output, please make sure that you get that output.

1. Construct adjacency matrix with the given data, assuming its graph and the weight of each edge is the rating given by user to the movie



the Adjacency matrix

$$\begin{matrix}
 & \begin{matrix} 1 & 2 & 3 \end{matrix} \\
 \begin{matrix} 1 \\ 2 \\ 3 \\ 4 \\ 5 \end{matrix} & \begin{bmatrix} 0 & 4 & 0 \\ 0 & 0 & 3 \\ 2 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 3 \end{bmatrix}
 \end{matrix}$$

you can construct this matrix like  $A[i][j] = r_{ij}$  here  $i$  is user\_id,  $j$  is movieid and  $r_{ij}$  is rating given by user  $i$  to the movie  $j$

Hint : you can create adjacency matrix using [csr\\_matrix](#)

1. We will Apply SVD decomposition on the Adjacency matrix [link1](#), [link2](#) and get three matrices  $U, \Sigma, V$  such that  $U \times \Sigma \times V^T = A$ ,  
if  $A$  is of dimensions  $N \times M$  then  
 $U$  is of  $N \times k$ ,  
 $\Sigma$  is of  $k \times k$  and  
 $V$  is  $M \times k$  dimensions.

\*. So the matrix  $U$  can be represented as matrix representation of users, where each row  $u_i$  represents a  $k$ -dimensional vector for a user

\*. So the matrix  $V$  can be represented as matrix representation of movies, where each row  $v_j$  represents a  $k$ -dimensional vector for a movie.

2. Compute  $\mu$ ,  $\mu$  represents the mean of all the rating given in the dataset.(write your code in [def m\\_u\(\)](#))
3. For each unique user initialize a bias value  $B_i$  to zero, so if we have  $N$  users  $B$  will be a  $N$  dimensional vector, the  $i^{th}$  value of the  $B$  will corresponds to the bias term for  $i^{th}$  user (write your code in [def initialize\(\)](#))

4. For each unique movie initialize a bias value  $C_j$  zero, so if we have  $M$  movies  $C$  will be a  $M$  dimensional vector, the  $j^{th}$  value of the  $C$  will corresponds to the bias term for  $j^{th}$  movie (write your code in `def initialize()`)
5. Compute  $dL/db\_i$  (Write you code in `def derivative_db()`)
6. Compute  $dL/dc\_j$ (write your code in `def derivative_dc()`)
7. Print the mean squared error with predicted ratings.

```

for each epoch:
    for each pair of (user, movie):
        b_i = b_i - learning_rate * dL/db_i
        c_j = c_j - learning_rate * dL/dc_j
    predict the ratings with formula

```

$$\hat{y}_{ij} = \mu + b_i + c_j + \text{dot\_product}(u_i, v_j)$$

1. you can choose any learning rate and regularization term in the range  $10^{-3}$  to  $10^2$
2. **bonus:** instead of using SVD decomposition you can learn the vectors  $u_i, v_j$  with the help of SGD algo similar to  $b_i$  and  $c_j$

```

In [1]: from google.colab import files
        files = files.upload()

```

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.  
Saving ratings\_train.csv to ratings\_train.csv

In [1]:

## Task 2

As we know  $U$  is the learned matrix of user vectors, with its  $i$ -th row as the vector  $u_i$  for user  $i$ . Each row of  $U$  can be seen as a "feature vector" for a particular user.

The question we'd like to investigate is this: do our computed per-user features that are optimized for predicting movie ratings contain anything to do with gender?

The provided data file [user\\_info.csv](#) contains an `is_male` column indicating which users in the dataset are male. Can you predict this signal given the features  $U$ ?

**Note 1 :** there is no train test split in the data, the goal of this assignment is to give an intuition about how to do matrix factorization with the help of SGD and application of truncated SVD. for better understanding of the collaborative filtering please check netflix case study.

**Note 2 :** Check if scaling of  $U, V$  matrices improve the metric

## Reading the csv file

```
In [2]: import pandas as pd
data=pd.read_csv('ratings_train.csv')
data.head()
```

```
Out[2]:
```

	user_id	item_id	rating
0	772	36	3
1	471	228	5
2	641	401	4
3	312	98	4
4	58	504	5

```
In [3]: data.shape
```

```
Out[3]: (89992, 3)
```

```
In [4]: data.describe()
```

```
Out[4]:
```

	user_id	item_id	rating
count	89992.000000	89992.000000	89992.000000
mean	461.579151	423.584663	3.529480
std	266.720677	330.264625	1.125686
min	0.000000	0.000000	1.000000
25%	253.000000	173.000000	3.000000
50%	446.000000	319.000000	4.000000
75%	681.000000	629.000000	4.000000
max	942.000000	1680.000000	5.000000

## Create your adjacency matrix

```
In [5]: from scipy.sparse import csr_matrix
adjacency_matrix = csr_matrix((data['rating'], (data['user_id'], data['item_id'])))
```

```
In [6]: adjacency_matrix.shape
```

```
Out[6]: (943, 1681)
```

## Grader function - 1

```
In [7]: def grader_matrix(matrix):
    assert(matrix.shape==(943,1681))
    return True
grader_matrix(adjacency_matrix)
```

Out[7]: True

**The unique items in the given csv file are 1662 only . But the id's vary from 0-1681 but they are not continuous and hence you'll get matrix of size 943x1681.**

### SVD decomposition

Sample code for SVD decomposition

```
In [8]: from sklearn.utils.extmath import randomized_svd
import numpy as np
matrix = np.random.random((20, 10))
U, Sigma, VT = randomized_svd(matrix, n_components=5, n_iter=5, random_state=None)
print(U.shape)
print(Sigma.shape)
print(VT.T.shape)

(20, 5)
(5,)
(10, 5)
```

### Write your code for SVD decomposition

```
In [27]: # Please use adjacency_matrix as matrix for SVD decomposition
# You can choose n_components as your choice
u, sigma, vt = randomized_svd(adjacency_matrix, n_components = 20, n_iter = 5, random_state=None)
print(u.shape)
print(sigma.shape)
print(vt.T.shape)

(943, 20)
(20,)
(1681, 20)
```

### Compute mean of ratings

```
In [10]: def m_u(ratings):
    '''In this function, we will compute mean for all the ratings'''
    # you can use mean() function to do this
    # check this (https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.mean.html)

    return ratings.mean()
```

```
In [11]: mu=m_u(data['rating'])
print(mu)

3.529480398257623
```

### Grader function -2

```
In [12]: def grader_mean(mu):
    assert(np.round(mu,3)==3.529)
    return True
mu=m_u(data['rating'])
grader_mean(mu)
```

Out[12]: True

### Initialize $B_i$ and $C_j$

Hint : Number of rows of adjacency matrix corresponds to user dimensions( $B_i$ ), number of columns of adjacency matrix corresponds to movie dimensions ( $C_j$ )

```
In [13]: def initialize(dim):  
    '''In this function, we will initialize bias value 'B' and 'C'.  
    # initialize the value to zeros  
    # return output as a list of zeros  
  
    return np.zeros(dim)
```

```
In [14]: dim = adjacency_matrix.shape[0] # give the number of dimensions for b_i (Here b_i corre  
b_i=initialize(dim)
```

```
In [15]: dim= adjacency_matrix.shape[1] # give the number of dimensions for c_j (Here c_j corresp  
c_j=initialize(dim)
```

### Grader function -3

```
In [16]: def grader_dim(b_i,c_j):  
    assert(len(b_i)==943 and np.sum(b_i)==0)  
    assert(len(c_j)==1681 and np.sum(c_j)==0)  
    return True  
grader_dim(b_i,c_j)
```

Out[16]: True

### Compute dL/db\_i

```
In [17]: def derivative_db(user_id,item_id,rating,U,V,mu,alpha):  
    '''In this function, we will compute dL/db_i'''  
    reg_term = 2*alpha*b_i[user_id]  
    loss_term = (-2)*(rating - mu - b_i[user_id] - c_j[item_id] - np.dot(U1[user_id],V1.  
  
    return reg_term+loss_term
```

### Grader function -4

```
In [18]: def grader_db(value):  
    assert(np.round(value,3)==-0.931)  
    return True  
U1, Sigma, V1 = randomized_svd(adjacency_matrix, n_components=2,n_iter=5, random_state=2  
# Please don't change random state  
# Here we are considering n_componets = 2 for our convinence  
alpha=0.01  
value=derivative_db(312,98,4,U1,V1,mu,alpha)  
grader_db(value)
```

Out[18]: True

### Compute dL/dc\_j

```
In [19]: def derivative_dc(user_id,item_id,rating,U,V,mu, alpha):  
    '''In this function, we will compute dL/dc_j'''  
    reg_term = 2*alpha*c_j[item_id]  
    loss_term = (-2)*(rating - mu - b_i[user_id] - c_j[item_id] - np.dot(U1[user_id],V1.  
  
    return reg_term+loss_term
```

```
In [20]: def grader_dc(value):
          assert(np.round(value,3)==-2.929)
          return True
U1, Sigma, V1 = randomized_svd(adjacency_matrix, n_components=2,n_iter=5, random_state=2
# Please don't change random state
# Here we are considering n_components = 2 for our convinence
alpha=0.01
value=derivative_dc(58,504,5,U1,V1,mu, alpha)
grader_dc(value)
```

Out[20]: True

### Compute MSE (mean squared error) for predicted ratings

for each epoch, print the MSE value

for each epoch:

for each pair of (user, movie):

$b_i = b_i - \text{learning\_rate} * dL/db_i$

$c_j = c_j - \text{learning\_rate} * dL/dc_j$

predict the ratings with formula

$$\hat{y}_{ij} = \mu + b_i + c_j + \text{dot\_product}(u_i, v_j)$$

```
In [51]: b_i = np.zeros(943)
          c_j = np.zeros(1681)
          from sklearn.metrics import mean_squared_error
          alpha = 0.5
          epochs = []
          error = []
          learning_rate = 0.0001
          for e in range(30):
              epochs.append(e+1)
              preds = []
              for user_id, item_id, rating in zip(data.iloc[:,0], data.iloc[:,1], data.iloc[:,2]):
                  db = derivative_db(user_id, item_id, rating, u, vt, mu, alpha)
                  b_i[user_id] = b_i[user_id] - learning_rate*db
                  dc = derivative_dc(user_id, item_id, rating, u, vt, mu, alpha)
                  c_j[item_id] = c_j[item_id] - learning_rate*dc
              for user_id, item_id, rating in zip(data.iloc[:,0], data.iloc[:,1], data.iloc[:,2]):
                  preds.append(mu+b_i[user_id]+c_j[item_id]+np.dot(u[user_id],vt.T[item_id]))
              error.append(mean_squared_error(data['rating'],np.array(preds)))
```


### Plot epoch number vs MSE

- epoch number on X-axis
- MSE on Y-axis

```
In [53]: import matplotlib.pyplot as plt

          plt.plot(epochs, error)
          plt.xlabel('Epoch Number')
```

```
plt.ylabel('Mean Squared Error')
plt.title('Epoch Number vs MSE')
plt.show()
```



## Task 2

- For this task you have to consider the user\_matrix U and the user\_info.csv file.
- You have to consider is\_male columns as output features and rest as input features. Now you have to fit a model by posing this problem as binary classification task.
- You can apply any model like Logistic regression or Decision tree and check the performance of the model.
- Do plot confusion matrix after fitting your model and write your observations how your model is performing in this task.
- Optional work- You can try scaling your U matrix. Scaling means changing the values of n\_components while performing svd and then check your results.

```
In [54]: from google.colab import files
files = files.upload()
```

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving user\_info.csv.txt to user\_info.csv.txt

```
In [55]: df = pd.read_csv('user_info.csv.txt')
df.head()
```

```
Out[55]:
```

	user_id	age	is_male	orig_user_id
--	---------	-----	---------	--------------

0	0	24	1	1
1	1	53	0	2
2	2	23	1	3
3	3	24	1	4
4	4	33	0	5




```
In [56]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix
import seaborn as sns

model = LogisticRegression(C = 0.01).fit(u,df['is_male'])
preds = model.predict(u)

print("Accuracy score: ", accuracy_score(df['is_male'], preds))
sns.heatmap(confusion_matrix(df['is_male'], preds), annot = True, cmap = 'Blues')
```

Accuracy score: 0.7104984093319194

Out[56]: <matplotlib.axes.\_subplots.AxesSubplot at 0x7fb5da87af10>



```
In [57]: df['is_male'].value_counts()
```

Out[57]:

1	670
0	273

Name: is\_male, dtype: int64

### OBSERVATIONS:

The user matrix we got on performing SVD doesn't have data to accommodate/ consider gender as a user\_info for recommendation tasks. The above logistic regression model seems to predict is\_male feature as 1 always. The model achieves even a 71% accuracy score due to the imbalanced nature of the user\_info dataset.

In [44]: